

## ER Diagram

- Many to Many

A can have multiple B and B can have multiple A



- One to Many

A can have multiple B but B can only have one A



- One to One

A can have only one B and B can have only one A



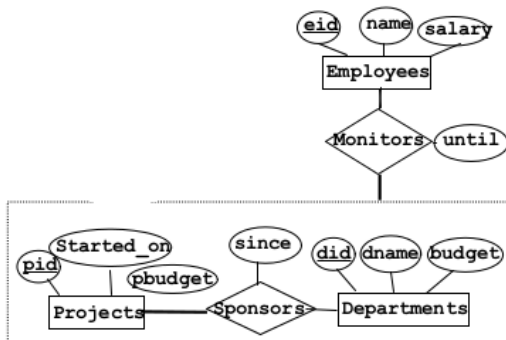
- At least one

Bold arrow specify there is at least one element.

	A have at least one
	A have exactly one

- Aggregation

Allows us to treat a relationship set R as an entity set so that R can participate in other relationships



## Relational algebra

Selection	$\sigma_{condition}(Element)$
Projection	$\pi_{attributes}(Element)$
Renaming	$\rho(R(A_1, \dots, A_n), R_{alias}(B_1, \dots, B_n))$
Cross product	$\times$
Join	$\bowtie$
Division	$\div$
Intersection	$\cap$
Union	$\cup$
Set different	$-$

- Condition/Theta Join  $R_{out} = R_{in1} \bowtie_{condition} R_{in2} = \sigma_{condition}(R_{in1} \times R_{in2})$
- Equi Join:  $R_{out} = R_{in1} \bowtie_{R_{in1}.a_1 = R_{in2}.b_1, \dots, R_{in1}.a_n = R_{in2}.b_n} R_{in2}$  Condition join where condition contains ONLY equalities
- Natural Join: Equijoin on all common attribute

## Sql

### Datatype

Char(n)	A character string of fixed length n
VarChar(n)	Denotes a string of up to n characters
INT or INTEGER	An integer
SHORTINT	Smaller integer
FLOAT or REAL	Float number
DOUBLE PRECISION	Double
DATE	Date format YYYY-MM-DD
TIME	Time format: hh:mm:ss

### Table operations

```
--Create table
CREATE TABLE Students
(
    id INT NOT NULL,
    name VARCHAR(20),
    login CHAR(10),
    major VARCHAR(20) DEFAULT 'undefined',
    school_id INT,
    PRIMARY KEY(id),
    FOREIGN KEY(school_id) REFERENCES School(id)
    CHECK (name != 'batman' OR name != 'joker')
)

--Drop table
DROP TABLE Students

--Alter table
ALTER TABLE Students ADD COLUMN firstyear:integer

--Creating index
CREATE INDEX useless_index ON table_name(column_name);
```

### Row operation

```
--INSERT
INSERT INTO Students (id, name, faculty) VALUES (8908998, 'Dupont', 'Science')

--Delete
DELETE FROM Students WHERE id = 0894984

--Update
UPDATE Students SET faculty = 'Arts' WHERE id = 9849849
```

### Trigger

```
CREATE TRIGGER updateSkater
AFTER DELETE ON Skaters
REFERENCING OLD TABLE AS DeletedSkaters
FOR EACH STATEMENT
INSERT
INTO StatisticsTable(ModTable, ModType, Count)
SELECT 'Skaters', 'delete', COUNT(*)
FROM DeletedSkaters

Use begin/end to encapsulate more than one
action
FOR EACH ROW/STATEMENT
WHEN ...
BEGIN ATOMIC
do 1thing;
do 2nd thing;
END
```

```

<!DOCTYPE DiscoverTheWorld [
<!ELEMENT DiscoverTheWorld (tour*,reservation*)>
<!ELEMENT tour (type, start-date, duration, price) >
<!ELEMENT reservation (cname, caddress, cost,
special*)>
<!ATTLIST tour TourId ID #REQUIRED >
<!ATTLIST reservation ResId ID #REQUIRED
TourId IDREF #REQUIRED>
<!ELEMENT type (#PCDATA) >
<!ELEMENT start-date (#PCDATA) >
<!ELEMENT duration (#PCDATA) >
<!ELEMENT price (#PCDATA) >
<!ELEMENT cname (#PCDATA) >
<!ELEMENT caddress (#PCDATA) >
<!ELEMENT cost (#PCDATA) >
<!ELEMENT special (#PCDATA) >
<!ATTLIST special price CDATA #REQUIRED>
]>

```

```

<!DOCTYPE Politics [ <!ELEMENT Politics (Politician*,
Province*)>
<!ELEMENT Politician ((CurrentMayor | CurrentMop)?,
address?)>
<!ELEMENT CurrentMayor (since?)>
<!ELEMENT CurrentMoP (since?)>
<!ELEMENT Province (City+, Riding+, population?)>
<!ELEMENT City (population?)>
<!ELEMENT Riding (population?)>
<!ELEMENT address (#PCDATA) >
<!ELEMENT since (#PCDATA) >
<!ELEMENT population (#PCDATA) >
<!ATTLIST Politician pname ID REQUIRED
website CDATA IMPLIED
friends IDREFS IMPLIED>
<!ATTLIST CurrentMayor cityID IDREF REQUIRED>
<!ATTLIST CurrentMoP rname IDREF REQUIRED>
<!ATTLIST Province pname ID REQUIRED>
<!ATTLIST City cityID ID REQUIRED
cname CDATA REQUIRED>
<!ATTLIST Riding rname CDATA REQUIRED>
]>

```

```

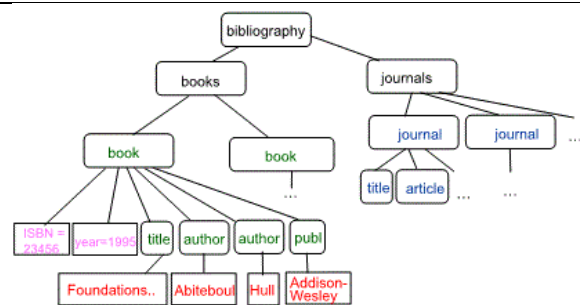
<bibliography>
  <books>
    <book ISBN="23456" year="1995">
      <title> Foundations ...
    </book>
    <book> ... </book>
  </books>
  <journals>
    <journal>
      <title> ... </title>
      <article> ... </article>
      ...
    </journal>
    <journal> ... </journal>
  </journals>
</bibliography>

```

```

<DiscoverTheWorld>
  <tour TourId="1">
    <type> Brazil jungle </type>
    <start-date> 16-April </start-date>
    <duration> 14 </duration>
    <price> 2229 </price>
  </tour>
  <tour TourId="2">
    <type> Brazil jungle </type>
    <start-date> 30-April </start-date>
    <duration> 21 </duration>
    <price> 2999 </price>
  </tour>
  <tour TourId="3">
    <type> Kenia safari </type>
    <start-date> 30-April </start-date>
    <duration> 21 </duration>
    <price> 3229 </price>
  </tour>
  <reservation ResId="541" TourId="1">
    <cname> Bettina Kemme </cname>
    <caddress> Montreal </caddress>
    <cost> 2579 </cost>
    <special price="5"> vegetarian </special>
    <special price="20"> single </special>
  </reservation>
  <reservation ResId="542" TourId="2">
    <cname> Your Name </cname>
    <caddress> Your Address </caddress>
    <cost> 3105 </cost>
    <special price="5"> vegetarian </special>
  </reservation>
</DiscoverTheWorld>

```



```

<!DOCTYPE people[
  <!ELEMENT people(person*)>
  <!ELEMENT person(name*, (lastname|familyname)?)>
  <!ATTLIST person PID ID #REQUIRED
    age CDATA #IMPLIED
    children IDREFS #IMPLIED
    mother IDREF #IMPLIED
  >
  <!ELEMENT name(#PCDATA)>
  <!ELEMENT lastname(#PCDATA)>
  <!ELEMENT familyname (#PCDATA)>
]>

```

Data types: PCDATA (parsed character data) or CDATA (unparsed)

#### Attributes

- ID unique identifier (similar to primary key)
- IDREF: reference to single ID
- IDREFS: space-separated list of references

#### Values

- can give a default value
- #REQUIRED must exist
- #IMPLIED optional

Specified in an XML file with <!DOCTYPE name SYSTEM "path/to/thing.dtd">

Can use regex style things too. \* is 0 or more. + is 1 or more, (a | b)? is one or the other

#### XPATH

- /bibliography/book/author all author elements by root navigating through those elements
- /bibliography/book/@ISBN All ISBN attributes
- //title all title elements anywhere in the document
- /bibliography/\*/title titles of bibliography entries assuming that there could be books, journals, reports, etc...
- /bibliography/book[@year>1995] returns books where the year > 1995
- /bibliography/book[author='FooBar']/@Year returns the years of books written by FooBar
- /bibliography/book[count(author) <2]
- /bibliography/book/author[position()=1]/name position is the location of the node in the node set

#### XQuery

where: condition

For	Let
for \$b in document("bib.xml")/bib/book return <result> \$b</result>	let \$b in document("bib.xml")/bib/book return <result> \$b</result>
<result><book>...</book></result> <result><book>...</book></result> <result><book>...</book></result> ... <result><book>...</book></result>	<result> <book>...</book> <book>...</book> .. <book>...</book> </result>



## Buffer

<p>DBMS stores information persistently on ( "hard" ) disks.</p> <p>□ Unit of transfer main-memory/disk: disk blocks or pages.</p> <p>□ Timing:</p> <ul style="list-style-type: none"> <li>☆ 2- 20 msec for random data block (bad seek time)</li> <li>☆ If blocks are sequentially on disk, only +1ms per block</li> <li>☆ Compare main memory access: in nanoseconds</li> </ul> <p>□ Basic operations (READ/WRITE from/to disk)</p> <p>□ Why disks?</p> <ul style="list-style-type: none"> <li>☆ Cheaper than Main Memory</li> <li>☆ Higher Capacity</li> <li>☆ Main Memory is volatile</li> </ul>	<p>When loading a page from disk:</p> <ul style="list-style-type: none"> <li>☆ Replacement frame must have "pin counter" of 0</li> </ul> <p>□ When requesting a page that is in the buffer</p> <ul style="list-style-type: none"> <li>☆ Increment pin counter</li> </ul> <p>□ After operation has finished</p> <ul style="list-style-type: none"> <li>☆ Decrement pin counter</li> <li>☆ Set dirty bit if page has been modified:</li> </ul> <p>□ Frame is chosen for replacement by a replacement policy:</p> <ul style="list-style-type: none"> <li>☆ Only unpinned page can be chosen (pin count = 0)</li> <li>☆ Least-recently-used (LRU), Clock, MRU etc.</li> </ul>	<p>If requested is not in pool:</p> <ul style="list-style-type: none"> <li>☆ If there is an empty frame, use it</li> <li>☆ Else choose an empty frame for replacement. If the frame is dirty (page was modified), write it to disk</li> <li>☆ Read requested page into chosen frame</li> </ul> <p>Buffer management in DBMS requires ability to:</p> <ul style="list-style-type: none"> <li>☆ <b>pin a page</b> in buffer pool, <b>force a page to disk</b> (important for implementing CC &amp; recovery),</li> <li>☆ adjust <b>replacement policy</b>, and <b>pre-fetch pages</b> based on access patterns in typical DB operations.</li> </ul>
--	---	---

## Indexing

COST MODEL	HEAP FILES	SORTED FILES
<p>Measure performances by simplifying the parameters (10 focused):</p> <ul style="list-style-type: none"> <li>☆ only consider disk reads (ignore writes)</li> <li>☆ only consider number of I/Os and not the individual time for each read (ignores page pre-fetch)</li> <li>☆ Average-case analysis: based on several simplistic assumptions. ●</li> </ul> <p>delete/update</p> <p>▲ depends on where</p>	<ul style="list-style-type: none"> <li>☆ Linked, unordered list of all pages of the file</li> <li>☆ Is it good for: <ul style="list-style-type: none"> <li>● scan retrieving all records (SELECT *)? <ul style="list-style-type: none"> <li>▲ yes, you have to retrieve all pages anyway</li> </ul> </li> <li>● equality search on primary key <ul style="list-style-type: none"> <li>▲ not great: have to read on avg half the pages for 1 record</li> </ul> </li> <li>● range search or equality search on non-primary key <ul style="list-style-type: none"> <li>▲ not great, all pages need to be read</li> </ul> </li> <li>● insert <ul style="list-style-type: none"> <li>▲ yes, can insert anywhere</li> </ul> </li> <li>● delete/update <ul style="list-style-type: none"> <li>▲ depends on where</li> </ul> </li> </ul> </li> </ul>	<ul style="list-style-type: none"> <li>☆ Records are ordered according to one or more attributes of the relation</li> <li>☆ Is it good for: <ul style="list-style-type: none"> <li>● scan retrieving all records (SELECT *)? <ul style="list-style-type: none"> <li>▲ yes, you have to retrieve all pages anyway</li> </ul> </li> <li>● equality search on sort attribute <ul style="list-style-type: none"> <li>▲ good: find first qualifying page with binary search (log2)</li> </ul> </li> <li>● range search on sort attribute <ul style="list-style-type: none"> <li>▲ good: find first qualifying page with binary search (log2): adjacent pages might have additional matching records</li> </ul> </li> </ul> </li> </ul>

Let suppose we have a relation R (A, B, C, D, F) such that:

- A and B are int (6 byte)
- C-F are char [40] (10 byte per char).
- Tuple = 172 bytes. 200,000 tuples
- Each data page has 4000 bytes and is around 80% full
- B values are uniformly distributed
- Rid = 10 bytes
- Size of pointer in intermediate page = 8 bytes
- Index pages are 4K and between 50%-100% full

Goal	Formula	With this example
Number of pages	$\frac{\text{number of tuples} * \text{tuple size}}{\text{fill rate} * \text{page size}}$	$\frac{172 * 200000}{40000 * 0.80} = 10750$
Index entry size in root and intermediate pages	$\text{size of key} + \text{size of pointer}$	$6 + 8 = 14 \text{ bytes}$
Average number of rids per data entry	$\frac{\text{number of tuples}}{\text{different values (if uniform)}}$	$\frac{200,000}{20,000} = 10$
Average length per data entry	$\text{size of key} + (\text{number of rids} * \text{size of rid})$	$6 + 10 * 10 = 106$
Average number of data entries per leaf page	$\frac{\text{fillrate} * \text{page size}}{\text{length of data entry}}$	$\frac{0.75 * 4000}{106} = 28 \text{ entries per page}$
Estimate number of leaf page	$\frac{\text{number of different values}}{\text{number of entrier per page}}$	$\frac{20,000}{28} = 715$
Number of entries in intermediate pages	$\frac{\text{fillrate} * \text{page size}}{\text{lenght of index enty}}$	$\min = \frac{0.5 * 4000}{14} = 143, \max = \frac{1 * 4000}{14} = 285$
Height of tree	$(\text{nb of entry in intermediate page})^{h-1} > \text{nb of leaf page}$	3

Non-clustered index B-tree with <k, list of rid>

Height of tree = Number of leaf pages / (min | max)? number of entries in intermediate pages

Give the pids of all projects within department D2 that started in 2014.

$$\pi_{pid} \left( \sigma_{dept_{id}=D2 \wedge start_{date}=2014}(Project) \right)$$

Give the pids of all projects that have at least one excellent evaluation

$$\pi_{Project.pid} \left( \sigma_{Evaluation.grade='excellent'}(Project \bowtie Evaluation) \right)$$

<b>Force Flush strategy</b> <ul style="list-style-type: none"> <li>All changes are flush to disk BEFORE commit</li> </ul>	<ul style="list-style-type: none"> <li>Completed transaction need not action</li> <li>Active transaction might have partial changes on disk(Need undone)</li> </ul>	<ul style="list-style-type: none"> <li>Append to log file log record before flushing</li> <li>At commit/abort append to log file commit/abort log record</li> <li>When recovering from crash: Scan log backward for each record if committed ignore otherwise install Before-Image of the record</li> </ul>
<b>No force flush strategy</b> <ul style="list-style-type: none"> <li>Changes might be flushed at any time(BEFORE/AFTER commit)</li> </ul>	<ul style="list-style-type: none"> <li>Done transaction might have missing changes (must be redone)</li> <li>Active/Aborted transaction might have been flushed before crash(Must be undone)</li> </ul>	<ul style="list-style-type: none"> <li>For each write(x) of a transaction T with x being on page P: Log record with before AND after image of x(Before so you can undo changes, After so you can redo changes)</li> <li>Flush before-image to disk before flushing the P</li> <li>Flush after-image to disk before commit of T</li> <li>At commit/abort append commit/abort record to log file and flush</li> </ul>

<ul style="list-style-type: none"> <li>• <b>Unrepeatable read:</b> If T1 read twice the same data item but T2 change its value between the first and the second</li> <li>• <b>Dirty read:</b> If T2 read from T1 before T1 commit.</li> <li>• <b>Lost update:</b> If T2 modify a data item modified by T1 without taking in account the value modified by T1.</li> </ul>	
--	--

<p>Schedule</p> <ul style="list-style-type: none"> <li>• <u>Serial schedule:</u> All transaction one after the other</li> <li>• <u>Non-serial schedule:</u> Transaction overlap <ul style="list-style-type: none"> <li>- <u>Serializable:</u> Dependency graph has no cycle(T1 always does action before T2)</li> <li>- <u>Recoverable schedule:</u> If transaction <math>T_i</math> reads a value written by transaction <math>T_j</math> then <math>T_i</math> commit only after <math>T_j</math> committed</li> <li>- <u>Avoiding cascading aborts:</u> A transaction reads only values written by committed transactions.</li> <li>- <u>Strict:</u> A transaction only read or overwrite value written by committed transaction</li> </ul> </li> </ul>	<p>Schedule examples:</p> <ul style="list-style-type: none"> <li>• Strict and serializable <math display="block">r1(x), w2(x), c2, w1(x), c1</math> </li> <li>• Avoids cascading aborts, non-strict, serializable</li> <li>• Recoverable, not avoiding cascade aborts, serializable <math display="block">r1(x), w2(y), w2(x), r1(y), c2, c1</math> </li> <li>• Not recoverable, serializable <math display="block">r1(x), w2(y), r1(x), c1, c2</math> </li> <li>• Not recoverable, Not-serializable</li> </ul>
--	---

Unrecoverable	Recoverable schedule with cascading abort	Recoverable schedule with commit	Avoids cascading	Non strict	Strict	Strict and serializable
<div>T1</div> <div>R(A)</div> <div>W(A)</div> <div></div> <div>R(A)</div> <div>commit</div> <div>commit</div>	<div>T1</div> <div>R(A)</div> <div>W(A)</div> <div></div> <div>R(A)</div> <div>abort</div> <div>abort</div>	<div>T1</div> <div>R(A)</div> <div>W(A)</div> <div></div> <div>R(A)</div> <div>commit</div> <div>commit</div>	<div>T1</div> <div>R(A)</div> <div>W(A)</div> <div>abort</div> <div></div> <div>R(A)</div> <div>commit</div>	<div>T1</div> <div>W(A)</div> <div></div> <div>W(A)</div> <div>abort</div> <div></div> <div>commit</div>	<div>T1</div> <div>W(A)</div> <div>abort</div> <div></div> <div>W(A)</div> <div>commit</div> <div>commit</div>	<div>T1</div> <div>R(x)</div> <div></div> <div></div> <div></div> <div></div> <div>W(x)</div> <div>commit</div> <div>W(x)</div> <div>commit</div>

<p><u>Lock request:</u></p> <ul style="list-style-type: none"> <li>• If lock is S, no X lock is active and the request queue is empty: <ul style="list-style-type: none"> <li>- Add the lock to the granted lock queue and set the lock type to S</li> </ul> </li> <li>• If lock is X and no lock active(request queue is also empty): <ul style="list-style-type: none"> <li>- Add the lock to the granted lock queue and set the lock type to X</li> </ul> </li> <li>• Otherwise <ul style="list-style-type: none"> <li>- Add the lock to the request lock queue</li> </ul> </li> </ul>	<p><u>Lock release:</u></p> <ul style="list-style-type: none"> <li>• Remove the lock from the granted lock queue</li> <li>• If this was the only lock granted on this object: <ul style="list-style-type: none"> <li>- Grant one X lock(If the first of the request is a X lock)</li> <li>- Grant n S lock(If the first n element are S lock)</li> </ul> </li> </ul>
<p><u>Deadlocks:</u></p> <ul style="list-style-type: none"> <li>• Make the wait-for graph(<math>T_i</math> need ressource lock by <math>T_j</math>)</li> <li>• If cycle then we have a deadlock (Noooooooooooooo...)</li> </ul>	<p><u>Solve Deadlock:</u></p> <ul style="list-style-type: none"> <li>• Add a timeout for each transaction and abort if transaction timeout. Problem on what timeout value to choose</li> <li>• Request all the lock at the beginning of the transaction</li> </ul>



**Predicate locking:**

- Grant lock on all records that satisfies logical predicates(e.g.  $depid > 5, age > 2 * salary$ )
- More bullshit

**Predicate locking example:**

- Assume 2 transactions:
  - *UPDATE Skaters set rating = 7 WHERE sid = 123*
  - *SELECT max(age) FROM Skaters WHERE rating = 5*
- Assume: T1 execute first then it has a X-lock on Skaters with sid=123
- Assume: T2 has to scan the entire table to get skater with rating=5
  - For each tuple
    - set S-lock on tuple
    - Check condition
    - If condition TRUE keep lock and return value
    - If condition FALSE release lock
  - It need to read the tuple where sid=123 and rating= 5 but block has T1 has a lock on it.
  - T2 is block by T1 although there is no conflict

**Problems of strict 2PL locking:**

- Very restrictive, low concurrency, problem with long query
- More and more exception

**In order to allow for more concurrency, SQL2 defines various levels of isolation**

- Assumed to be implemented by different forms of locking
- Avoid different levels of anomalies
- Used for non-critical transactions or read-only transactions
- Lower levels of isolation do NOT provide serializability

**Problem**

- Definitions are no more appropriate if systems do not use locking but other forms of concurrency control
- For instance, Oracle's "serializable" level does not provide serializable schedule as defined in the literature

**Isolation level:**

- In principle isolation levels are independent of concurrency control mechanics
- In reality they were defined with locking in mind

Isolation level/Anomaly	Dirty read	Unrepeatable read	Phantom
Read uncommitted	Maybe	Maybe	Maybe
Read committed	No	Maybe	Maybe
Repeated reads	No	No	Maybe
Serializable	No	No	No

- **Read uncommitted:**
  - Read op. do not set locks; can read not-committed updates
- **Read committed:**
  - Read op. set short S locks; have to wait for X locks to be released
  - release lock immediately after execution of op
- **Repeated reads:**
  - Read operations set standard lock S locks; standard 2PL
- **Serializable:**
  - Read op. must set S locks that cover all objects that are read
  - predicate locks or coarse locks (e.g., lock on entire relation)

# Big data

Some bullshit info:

- Hardware
  - CPU does not increase
  - Instead muticode
- Usage
  - Astronomy: high-resolution, high-frequency sky surveys
  - Medicine: digital records, MRI, ultrasound
  - Biology: sequencing data
  - User behavior data: click streams, search logs

## Horizontal data Partitioning:

- Data
  - Large table  $R(K, A, B, C)$
  - Key value store  $KV(K, V)$
- Goal
  - Partition into chunks  $c_1, \dots, c_n$  of records stored at N nodes
- Range partition
  - Equal size of each chunk
- Hash partitioned on attribute X
  - Record  $r$  goes to chunk  $i$ , according to hash function
  - xample: hash function  $H(r.X) \bmod P+1$
- Range partitioned on attribute X
  - Partition range of X into:  $-\infty = v_0 < v_1 < \dots < v_p = \infty$
  - Record  $r$  goes to chunk  $i$ , if  $v_{i-1} < t.X < v_i$

## Execution steps:


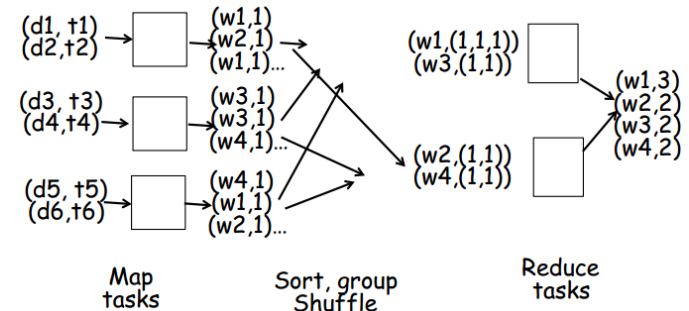
- User indicates  $m$  (number of map tasks),  $r$  (number of reduce tasks),  
key/value set = document Set DS
- System creates  $m$  map tasks and splits input set of 1key/value pairs into  $m$  partitions and gives each map task one partition as input
- Each Map task executes user written map function
  - WordCountMap:
    - For each input key/value pair  $(dkey, dtext)$
    - For each word  $w$  of  $dtext$
    - Output key-value pair  $(w, 1)$
- Next step only completes once all map tasks have completed
- System sorts map outputs by key and transforms all key/value pairs  $(k, v_1), (k, v_2), \dots, (k, v_n)$  with same key  $k$  to one key/value-list pair  $(k, (v_1, v_2, \dots, v_n))$ 
  - For Word count: all  $(\text{'and'}, 1), (\text{'and'}, 1), (\text{'and'}, 1) \dots$  are transformed into one  $(\text{'and'}, (1, 1, 1, \dots))$
- System partitions output by key into  $r$  partitions and assigns these partitions as inputs to the  $r$  reduce tasks
- Each reduce task executes user written reduce function
  - WordCountReduce:
    - For each input key/value-list pair  $(k, (v_1, v_2, \dots, v_n))$
    - Output  $(k, n)$

Parallel Query Evaluation:


- Inter-query parallelism
  - Different queries run in parallel on different processors; each query is executed sequentially
- Inter-operator parallelism
  - Different operator within the same execution tree run on different processors
- Intra-operator parallelism
  - A single operator(JOIN, GROUP, ...) runs on many processor

Vertical Partitioning:

- Column stores
- Data: relation  $R(K, A, B, C)$
- Partition into  $RA(K, A), RB(K, B), RC(K, C)$
- Query:
  - SELECT A FROM R
- Query only needs to access partition RA
- Much less IO



## Map reduce

<p><b>Relational Operators with Map/ reduce</b></p> <ul style="list-style-type: none"> <li>Assume <math>R(A, B, C)</math> relation (no duplicates)</li> <li><b>Selection with condition <math>c</math> on <math>R</math></b> <ul style="list-style-type: none"> <li>for each tuple <math>t</math> of <math>R</math> for which condition <math>c</math> holds, output <math>(t, t)</math></li> <li>Reduce: identity, that is output <math>(t, t)</math></li> </ul> </li> <li><b>Projection on <math>A, B</math>, of <math>R</math></b> <ul style="list-style-type: none"> <li><u>Map</u>: transform each tuple <math>t = (a, b, c)</math> of <math>R</math> into tuple <math>t' = (a, b)</math> of <math>R</math>, and output <math>(t', t')</math></li> <li>There might now be duplicates, that is several <math>(t', t')</math> tuples, the group function will aggregate them to <math>(t', (t', \dots, t'))</math></li> <li><u>Reduce</u>: for each tuple <math>(t', (t', \dots, t'))</math> output <math>(t', t')</math></li> </ul> </li> </ul>	<ul style="list-style-type: none"> <li><b>Grouping: <math>\text{SELECT } a, \text{sum}(b) \text{ GROUP by } (a)</math></b> <ul style="list-style-type: none"> <li><u>Map</u>: for each tuple <math>(a, b, c)</math> of <math>R</math> output <math>(a, b)</math></li> <li>Group and shuffle will create for each value <math>a</math> a key/value-list <math>(a, (b_1, b_2, \dots))</math></li> <li><u>Reduce</u>: for each <math>(a, (b_1, b_2, \dots))</math> perform aggregation (e.g., <math>b_1 + b_2, \dots</math>)</li> </ul> </li> </ul>
<ul style="list-style-type: none"> <li><b>Natural Join <math>R(A, B, C)</math> with <math>Q(C, D, E)</math> via hash-join(<math>\text{SELECT FROM } R_1, R_2</math>)</b> <ul style="list-style-type: none"> <li><u>Map</u>:           <ul style="list-style-type: none"> <li>For each tuple <math>(a, b, c)</math> of <math>R</math>, output <math>(c, (R, (a, b)))</math></li> <li>For each tuple <math>(c, d, e)</math> of <math>Q</math>, output <math>(c, (Q, (d, e)))</math></li> </ul> </li> <li>Group and shuffle will aggregate all key/value pairs with same <math>c</math>-value</li> <li><u>Reduce</u>:           <ul style="list-style-type: none"> <li>For each tuple <math>(c, \text{value-list})</math>, example:               <ul style="list-style-type: none"> <li><math>(\text{value-list} = (R, (a_1, b_1)), (R, (a_2, b_2)), \dots (Q, (d_1, e_1)), \dots)</math></li> <li><math>R_t = Q_t = \text{empty}</math>;</li> <li>for each <math>v = (\text{rel}, \text{tuple})</math> in value-list</li> <li>if <math>v.\text{rel} = R</math>: insert tuple into <math>R_t</math> else insert tuple into <math>Q_t</math></li> <li>for <math>v_1</math> in <math>R_t</math>, for <math>v_2</math> in <math>Q_t</math>, output <math>(c, v_1, v_2)</math></li> </ul> </li> <li>Basically produces all combinations <math>(c, a_i, b_j, d_k, e_l)</math></li> </ul> </li> </ul> </li> </ul>	<p>Pig latin:</p> <ul style="list-style-type: none"> <li>Users = <b>LOAD</b> 'users' <b>AS</b> (name,age);</li> <li>Filtered = <b>FILTER</b> Users <b>BY</b> age &gt;= 18 <b>AND</b> age &lt;= 25;</li> <li>Pages = <b>LOAD</b> 'pages' <b>AS</b> (uname, url);</li> <li>Joined = <b>JOIN</b> Fltrd <b>BY</b> name, Pages <b>BY</b> uname;</li> <li>Grouped = <b>GROUP</b> Jnd <b>BY</b> url;</li> <li>Smmmd = <b>FOREACH</b> Grpd <b>GENERATE</b> (\$0), COUNT(\$1) <b>AS</b> clicks;</li> <li>Srtd = <b>ORDER</b> Smmmd <b>BY</b> clicks desc;</li> <li>Top5 = <b>LIMIT</b> Srtd 5;</li> <li><b>STORE</b> Top5 <b>INTO</b> 'top5sites'</li> </ul> 
<p><b>Pig examples:</b></p> <pre>raw = LOAD ....  -- filter per percent fltrd = FILTER raw by percent &gt;= 60;  gen = foreach fltrd generate CONCAT(firstname, CONCAT(' ', lastname)); results = DISTINCT gen; STORE results INTO 's3n://comp421-h4/q1_results';</pre>	<pre>raw = LOAD .....  --some data entries use the middle name as well, so this way we will catch all of them fltrd = FILTER raw by votes &gt;= 100; SPLIT fltrd INTO winners IF elected == 1, losers IF elected == 0;  elections = JOIN winners BY (date, type, parl, prov, riding), losers BY (date, type, parl, prov, riding);  vote_differences = foreach elections generate winners::lastname as winner, losers::lastname as loser, (winners::votes-losers::votes) as vote_difference:int;  results = FILTER vote_differences by vote_difference &lt; 10; --print the result tuple to the screen  STORE results INTO 's3n://comp421-h4/q2_results';</pre>

```

raw = LOAD ...
--some data entries use the middle name as well, so this way
we will catch all of them
fltrd = FILTER raw by type == 'Gen' and elected == 1;

parl_group = GROUP fltrd BY parl;

parl_count = FOREACH parl_group GENERATE ($0) as parl,
COUNT($1) as count;
parl_count_before = FOREACH parl_count GENERATE ($0+1) as
parl, $1 as count;
parl_join = JOIN parl_count BY parl, parl_count_before BY
parl;

parl_diff = FOREACH parl_join GENERATE parl_count::parl as
parl, parl_count::count, parl_count::count -
parl_count_before::count;

results = ORDER parl_diff BY parl;

dump results;

```

```

raw = LOAD ...
parl_group = GROUP raw BY parl;

parl_count = FOREACH parl_group GENERATE ($0) as parl,
COUNT($1) as parl_count;

party_group = GROUP raw BY (parl, party);

party_count = FOREACH party_group GENERATE FLATTEN($0) as
(parl, party), COUNT($1) as party_count;

joined = JOIN party_count BY parl, parl_count BY parl;

results = FOREACH joined GENERATE parl_count::parl as parl,
party_count::party as party, party_count::party_count as
party_count, parl_count::parl_count as parl_count;

store results into '/user/hadoop/q4output.csv' using
PigStorage('\t', '-schema');

```

## Query evaluation

<p><b>Examples for flowing problems:</b></p> <ul style="list-style-type: none"> <li>Participates <ul style="list-style-type: none"> <li>100 000 tuples</li> <li>1000 pages</li> <li>100 tuples per page</li> </ul> </li> <li>Skaters <ul style="list-style-type: none"> <li>40 000 tuples</li> <li>500 pages</li> <li>80 tuples per page</li> <li>Index on sid has 170 leaf page</li> <li>Index on names has 300 leaf page</li> </ul> </li> </ul>	<ul style="list-style-type: none"> <li><b>Reduction factor of a condition defined as</b> <ul style="list-style-type: none"> <li><math>Red(\sigma_{condition}(R)) =  \sigma_{condition}(R) / R </math></li> <li><math>Red(\sigma_{rating=5}(Skaters)) = \frac{ \sigma_{rating=5}(Skaters) }{ Skaters } = \frac{15}{100} = 0.15</math></li> </ul> </li> <li><b>If not known, DBMS makes simple assumptions</b> <ul style="list-style-type: none"> <li><math>Red(\sigma_{rating=5}(Skaters)) = \frac{1}{ diff\ rating } = 0.1</math></li> <li><math>Red(\sigma_{age \leq 10}(Skaters)) = \frac{10 - \min(age)}{\max(age) - \min(age)} = \frac{10 - 4}{30 - 4} = \frac{6}{26} = 0.025</math></li> </ul> </li> <li><b>Result size: number of input tuples * reduction factor</b></li> <li>How to know the number of different values, max, min, <ul style="list-style-type: none"> <li>through indices, heuristics, separate statistics (histograms)</li> </ul> </li> </ul>
<p><b>Simple selection:</b></p> <ul style="list-style-type: none"> <li>No index: <ul style="list-style-type: none"> <li>Search on arbitrary attributes: scan the entire relation e.g. <math>cost = page(Skaters)</math></li> <li>Search on primary key attributes: scan on average half of S e.g. <math>cost = \frac{page(Skaters)}{2} = 250</math></li> </ul> </li> <li>Index on selection attribute <ul style="list-style-type: none"> <li>Use index to find qualifying data entries, then retrieve corresponding data records.</li> </ul> </li> </ul>	<p><b>Clustered B+tree</b></p> <ul style="list-style-type: none"> <li>Costs: <ul style="list-style-type: none"> <li><math>\#LeafPages + \#dataPageToRetrieve</math></li> </ul> </li> <li>Example 1 <ul style="list-style-type: none"> <li>SELECT * FROM Skaters WHERE name sid = 5</li> <li>1 tuple match</li> <li><math>cost = 1\ leaf\ page + 1\ data\ page = 2</math></li> </ul> </li> <li>Example 2 <ul style="list-style-type: none"> <li>SELECT * FROM Skaters WHERE name LIKE 'Z%'</li> <li>System estimate the number of matching tuples(Around 100 match on 2 data page as its clustered)</li> <li><math>cost = 1\ leaf\ page + 2\ data\ page = 3</math></li> </ul> </li> <li>Example 3 <ul style="list-style-type: none"> <li>SELECT * FROM Skaters WHERE name &lt; 'F%'</li> <li>Assume around 10 000 tuples match(On 125 datapage)</li> <li><math>cost = 1\ leaf\ page + 125\ data\ pages = 126</math></li> </ul> </li> </ul>
<p><b>Unclustered B+tree</b></p> <ul style="list-style-type: none"> <li>Costs: <ul style="list-style-type: none"> <li><math>\#dataPageToRetrieve = \#tuples</math></li> <li><math>\#LeafPages + \#dataPageToRetrieve</math></li> </ul> </li> <li>Example 1 <ul style="list-style-type: none"> <li>SELECT * FROM Skaters WHERE name sid = 5</li> <li>Same</li> </ul> </li> <li>Example 2 <ul style="list-style-type: none"> <li>SELECT * FROM Skaters WHERE name LIKE 'Z%'</li> <li>Assume around 100 match on 80 data page but we need to retrieve some data page twice</li> <li><math>cost = 1\ leaf\ page + 101\ data\ page = 101</math></li> </ul> </li> <li>Example 3 <ul style="list-style-type: none"> <li>SELECT * FROM Skaters WHERE name &lt; 'F%'</li> <li>Assume around 10 000 tuples match</li> <li><math>cost = 75\ leaf\ page + 10\ 000\ data\ pages = 10\ 001</math></li> </ul> </li> </ul>	<p><b>Unclustered B+tree with sorting:</b></p> <ul style="list-style-type: none"> <li>Sort matching data entries (rid=pid,slot-id) in leaf-pages by page-id</li> <li>Only fast if the 75 leaf pages with matching entries fit in main memory</li> <li>Retrieve each page only once and get all matching tuples</li> <li><math>\#data\ pages = \#data\ pages\ that\ have\ at\ least\ one\ matching\ tuple;</math></li> <li>worst case is total # of data pages</li> <li>Example 1 <ul style="list-style-type: none"> <li>SELECT * FROM Skaters WHERE name sid = 5</li> <li>Same</li> </ul> </li> <li>Example 2 <ul style="list-style-type: none"> <li>SELECT * FROM Skaters WHERE name LIKE 'Z%'</li> <li>Assume around 100 match on 80 data page</li> <li><math>cost = 1\ leaf\ page + 80\ data\ page = 81</math></li> </ul> </li> <li>Example 3 <ul style="list-style-type: none"> <li>SELECT * FROM Skaters WHERE name &lt; 'F%'</li> <li>Assume around 10 000 tuples match(assume thataround 490 data pages)</li> <li><math>cost = 75\ leaf\ page + 490\ data\ pages = 565</math></li> </ul> </li> <li>Note: sorting expensive if leaf-pages do not fit in main-memory</li> </ul>
<p><b>Sort:</b></p> <ul style="list-style-type: none"> <li>Sometimes a pass 2 is needed <ul style="list-style-type: none"> <li>Pass 0 created more runs than there are main memory buffers</li> <li>Therefore Pass 1 produces more than one run</li> <li>Pass 2 takes the runs of Pass 1 and merges them</li> </ul> </li> <li>Cost <ul style="list-style-type: none"> <li>SELECT sname, age FROM Skaters ORDER BY age</li> <li>If everything fits into main memory (Only pass 0 needed): <ul style="list-style-type: none"> <li>Read number of data pages</li> <li>sort and pipeline result into next operator (project)</li> </ul> </li> <li>Pass 0 + pass 1 needed <ul style="list-style-type: none"> <li>Pass 0: read # pages, write # pages (have to write temp. results!)</li> <li>Pass 1: read # pages, sort and pipeline result into next operator</li> <li>3 * #pages</li> </ul> </li> <li>Pass 0 + pass1 + pass2 needed <ul style="list-style-type: none"> <li>5 * #pages</li> </ul> </li> </ul> </li> </ul>	<p><b>Join cost estimation:</b></p> <ul style="list-style-type: none"> <li><math> Skaters Participates  =  Participates </math> <ul style="list-style-type: none"> <li>Join attribute is primary key for Skaters</li> <li>Each Participates tuple matches exactly with one Skaters tuple</li> </ul> </li> <li><math> Skaters \times Participates  =  Participates  *  Skaters </math> <ul style="list-style-type: none"> <li>Cross product is always the product of individual relation sizes</li> </ul> </li> <li>For other joins more difficult to estimate (Continues in next episode...)</li> </ul>

#### Nested loop joins:

- Simple nested loop join: For each tuple in the outer relation P, we scan the entire inner relation S
  - $cost = PartPages + CARD(P) * SkaterPages = 1000 + 100000 * 500$
- Page-oriented Nested Loops join: For each page of P, get each page of S, and write out matching pairs of tuples <p, s>, where p is in P-page and s is in S-page.
  - $cost = PartPage + PartPage * SkaterPage = 1000 + 1000 * 500$

#### Join cost on relation R1(outer) and R2(inner):

##### **Block oriented nested loop join:**

- Smaller relation fits in main memory+2extra buffer page:
$$cost = page(R1) + page(R2)$$
- No relation fits in main memory(B join frame):
$$cost = page(R1) + \frac{page(R2) * page(R1)}{B - 2}$$

##### **Index nested loop join**

- Index on the join column of one of the relation(R2):
$$cost = page(R1) + card(R1) * cost\_finding\_index(R2)$$
- If the join attribute is primary key in inner relation

##### **Sort merge join**

- Sort P and S on the join column, then scan them to do a "merge" (on join col.), and output result tuples
  - Advance scan of P until current P-tuple  $\geq$  current S tuple, then advance scan of S until current S-tuple  $\geq$  current P tuple; do this until current P tuple = current S tuple.
  - At this point, all P tuples with same value in  $P_i$  (current P group) and all S tuples with same value in  $S_j$  (current S group) match; output <p, s> for all pairs of such tuple 1s.
- P is scanned once; each S group is scanned once per matching P tuple. (Multiple scans of an S group are likely to find needed pages in buffer.)



© Original Artist / Search ID: shr1169

Rights Available from CartoonStock.com

"This problem had been my life's work. I planned to devote my remaining years to it. It's just been solved in four seconds."

#### THE #1 PROGRAMMER EXCUSE FOR LEGITIMATELY SLACKING OFF:

"MY CODE'S COMPILING."

